

## Introductory Applied Machine Learning

### Nearest Neighbour Methods


Victor Lavrenko and Nigel Goddard  
School of Informatics

### Overview

- Nearest neighbour method
  - classification and regression
  - practical issues: k, distance, ties, missing values
  - optimality and assumptions
- Making kNN fast:
  - K-D trees
  - inverted indices
  - fingerprinting
- References: W&F sections 4.7 and 6.4

Copyright © 2014 Victor Lavrenko

### Intuition for kNN




- set of points  $(x, y)$ 
  - two classes
- is the box red or blue
- how did you do it
  - use Bayes rule?
  - a decision tree?
  - fit a hyperplane?
- nearby points are red
  - use this as a basis for a learning algorithm

Copyright © 2014 Victor Lavrenko

### Nearest-neighbor classification


- Use the intuition to classify a new point  $x$ :
  - find the most similar training example  $x'$
  - predict its class  $y'$
- Voronoi tessellation
  - partitions space into regions
  - boundary: points at same distance from two different training examples
- classification boundary
  - non-linear, reflects classes well
  - compare to NB, DT, logistic
  - impressive for simple method



Copyright © 2014 Victor Lavrenko

### Nearest neighbour: outliers


- Algorithm is sensitive to outliers
  - single mislabeled example dramatically changes boundary
- No confidence  $P(y|x)$
- Insensitive to class prior
- Idea:
  - use more than one nearest neighbor to make decision
  - count class labels in  $k$  most similar training examples
    - many "triangles" will outweigh single "circle" outlier



Copyright © 2014 Victor Lavrenko

### Example: handwritten digits

- 16x16 bitmaps
- 8-bit grayscale
- Euclidian distance
  - over raw pixels
- Accuracy:
  - 7-NN ~ 95.2%
  - SVM ~ 95.8%
  - humans ~ 97.5%



Copyright © 2014 Victor Lavrenko

Figure © Sam Roweis, 2006

### kNN classification algorithm

- Given:
  - training examples  $\{x_p, y_p\}$ 
    - $x_i$  ... attribute-value representation of examples
    - $y_i$  ... class label: {ham, spam}, digit {0,1,...,9} etc.
  - testing point  $x$  that we want to classify
- Algorithm:
  - compute distance  $D(x, x_i)$  to every training example  $x_i$
  - select  $k$  closest instances  $x_{i_1} \dots x_{i_k}$  and their labels  $y_{i_1} \dots y_{i_k}$
  - output the class  $y^*$  which is most frequent in  $y_{i_1} \dots y_{i_k}$


Copyright © 2014 Victor Lavrenko

### kNN regression algorithm

- Given:
  - training examples  $\{x_p, y_p\}$ 
    - $x_i$  ... attribute-value representation of examples
    - $y_i$  ... real-valued target (profit, rating on YouTube, etc)
  - testing point  $x$  that we want to predict the target
- Algorithm:
  - compute distance  $D(x, x_i)$  to every training example  $x_i$
  - select  $k$  closest instances  $x_{i_1} \dots x_{i_k}$  and their labels  $y_{i_1} \dots y_{i_k}$
  - output the mean of  $y_{i_1} \dots y_{i_k}$ :
 
$$\hat{y} = f(x) = \frac{1}{k} \sum_{j=1}^k y_{i_j}$$


Copyright © 2014 Victor Lavrenko

## Example: kNN regression in 1-d



## Choosing the value of k

- Value of k has strong effect on kNN performance
  - large value  $\rightarrow$  everything classified as the most probable class:  $P(y)$
  - small value  $\rightarrow$  highly variable, unstable decision boundaries
    - small changes to training set  $\rightarrow$  large changes in classification
  - affects "smoothness" of the boundary
- Selecting the value of k
  - set aside a portion of the training data (validation set)
  - vary k, observe training  $\rightarrow$  validation error
  - pick k that gives best generalization performance

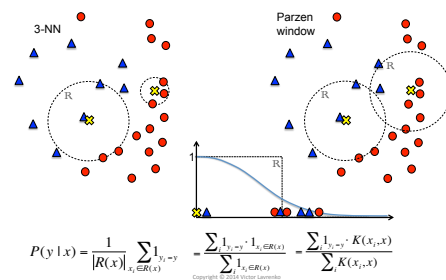


## kNN: practical issues

- Resolving ties:
  - equal number of positive/negative neighbours
  - use odd k (doesn't solve multi-class)
  - breaking ties:
    - random: flip a coin to decide positive / negative
    - prior: pick class with greater prior
    - nearest: use 1-NN classifier to decide
- Missing values
  - have to "fill in", otherwise can't compute distance
  - key concern: should affect distance as little as possible
  - reasonable choice: average value across entire dataset

Copyright © 2014 Victor Lumsden

## kNN, Parzen Windows and Kernels



Copyright © 2014 Victor Lumsden

## Distance measures

- Key component of the kNN algorithm
  - defines which examples are similar & which aren't
  - can have strong effect on performance
- Euclidian (numeric attributes):  $D(x, x') = \sqrt{\sum_d |x_d - x'_d|^2}$ 
  - symmetric, spherical, treats all dimensions equally
  - sensitive to extreme differences in single attribute
    - behaves like a "soft" logical OR
- Hamming (categorical attributes):  $D(x, x') = \sum_d 1_{x_d \neq x'_d}$ 
  - number of attributes where  $x, x'$  differ

Copyright © 2014 Victor Lumsden

## Distance measures (2)

- Minkowski distance ( $p$ -norm):  $D(x, x') = \sqrt[p]{\sum_d |x_d - x'_d|^p}$ 
  - $p=2$ : Euclidian
  - $p=1$ : Manhattan
  - $p=0$ : Hamming ... logical AND
  - $p=\infty$ :  $\max_d |x_d - x'_d|$  ... logical OR
- Kullback-Leibler (KL) divergence:
  - for histograms ( $x_d > 0, \sum_d x_d = 1$ ):  $D(x, x') = -\sum_d x_d \log \frac{x_d}{x'_d}$
  - asymmetric, excess bits to encode  $x$  with  $x'$
- Custom distance measures (*BM25* for text)

Copyright © 2014 Victor Lumsden

## kNN pros and cons

- Almost no assumptions about the data
  - smoothness: nearby regions of space  $\rightarrow$  same class
  - assumptions implied by distance function (only locally!)
  - non-parametric approach: "let the data speak for itself"
    - nothing to infer from the data, except  $k$  and possibly  $D()$
    - easy to update in online setting: just add new item to training set
- Need to handle missing data: fill-in or create a special distance
- Sensitive to class-outliers (mislabelled training instances)
- Sensitive to lots of irrelevant attributes (affect distance)
- Computationally expensive:
  - space: need to store all training examples
  - time: need to compute distance to all examples:  $O(nd)$ 
    - $n$  ... number of training examples,  $d$  ... cost of computing distance
    - $n$  grows  $\rightarrow$  system will become slower and slower
    - expense is at testing, not training time (bad)

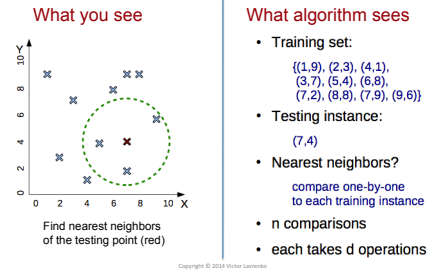
Copyright © 2014 Victor Lumsden

## Summary: kNN

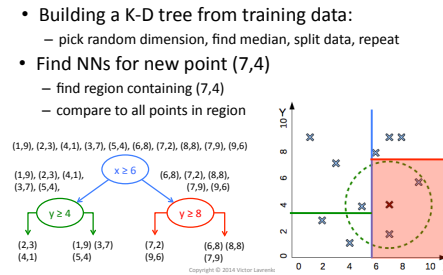
- Key idea: nearby points  $\rightarrow$  same class
  - important to select good distance function
- Can be used for classification and regression
- Simple, non-linear, asymptotically optimal
  - does not make assumptions about the data
  - "let the data speak for itself"
- Select  $k$  by optimizing error on held-out set
- Naïve implementations slow for big datasets
  - use K-D trees (low-d) or inverted lists (high-d)

Copyright © 2014 Victor Lumsden

## Why is kNN slow?



## K-D tree example



## Making kNN fast

- Training:  $O(d)$ , but testing:  $O(nd)$
  - Reduce  $d$ : dimensionality reduction
    - simple feature selection, other methods  $O(d^3)$
  - Reduce  $n$ : don't compare to **all** training examples
    - idea: quickly identify  $m \ll n$  potential near neighbors
      - compare only to those, pick  $k$  nearest neighbors  $\rightarrow O(md)$  time
    - K-D trees**: low-dimensional, real-valued data
      - $O(d \log_2 n)$ , only works when  $d \ll n$ , inexact: may miss neighbors
    - inverted lists**: high-dimensional, discrete data
      - $O(n'd)$  where  $d' \ll d$ ,  $n' \ll n$ , only for sparse data (e.g. text), exact
    - locality-sensitive hashing**: high-d, discrete or real-valued
      - $O(n'd)$ ,  $n' \ll n$  ... bits in fingerprint, inexact: may miss near neighbors
- Copyright © 2014 Victor Lumsden

## Locality-Sensitive Hashing (LSH)

- Random hyper-planes  $h_1 \dots h_k$ 
    - space sliced into  $2^k$  regions (polytopes)
    - compare  $x$  only to training points in the same region  $R$
  - Complexity:  $O(kd + dn/2^k)$ 
    - $O(kd)$  to find region  $R$ ,  $k \ll n$ 
      - dot-product  $x$  with  $h_1 \dots h_k$
    - compare to  $n/2^k$  points in  $R$
  - Inexact: missed neighbors
    - repeat with different  $h_1 \dots h_k$
  - Why not K-D tree?
- 
- Copyright © 2014 Victor Lumsden

## Inverted list example

- Data structure used by search engines (Google, etc)
    - list all training examples that contain particular attribute
    - assumption: most attribute values are zero (sparseness)
  - Given a new testing example:
    - merge inverted lists for attributes present in new example
    - $O(dn)$ :  $d$  ... nonzero attributes,  $n$  ... avg. length of inverted list
- |                             |      |  |            |          |          |          |          |          |
|-----------------------------|------|--|------------|----------|----------|----------|----------|----------|
| D1: "send your password"    | spam |  | send →     | <u>1</u> | <u>2</u> | <u>3</u> | <u>4</u> | <u>5</u> |
| D2: "send us review"        | ham  |  |            |          |          |          |          |          |
| D3: "send us password"      | spam |  | your →     | <u>1</u> | <u>5</u> | <u>6</u> |          |          |
| D4: "send us details"       | ham  |  | review →   | <u>2</u> | <u>6</u> |          |          |          |
| D5: "send your password"    | spam |  | account →  |          | <u>6</u> |          |          |          |
| D6: "review your account"   | ham  |  | password → | <u>1</u> | <u>3</u> | <u>5</u> |          |          |
| new email: "account review" |      |  |            |          |          |          |          |          |
- Copyright © 2014 Victor Lumsden